

Lakehouse Architecture Basics

Day 3 - Databricks 14-Day AI Challenge

Yash Kavaiya

Gen AI Guru | Easy AI Labs

January 11, 2026

Agenda

- **Data Architectures Evolution** - From Warehouses to Lakehouse
- **What is a Data Lakehouse?** - Combining the best of both worlds
- **Delta Lake** - The foundation of Lakehouse
- **Lakehouse Components** - Complete architecture overview
- **ACID Transactions** - Ensuring data reliability
- **Medallion Architecture** - Bronze, Silver, Gold layers

The Journey to Modern Data Architecture


Phase 1: Data Warehouses (1980s-2000s)

- Structured data only
- ACID transactions
- SQL-based querying
- High cost per TB

Phase 2: Data Lakes (2010s)


- All data types
- Schema-on-read
- Low cost storage
- Poor reliability

Phase 1: Data Warehouses



Limitation: Not suitable for unstructured data (images, logs, JSON)

Phase 2: Data Lakes



Problem: No ACID transactions → “**Data Swamps**”

The Two-Tier Architecture Problem




Problems:

- • Data duplication
- • Stale data
- • Complex pipelines
- • High costs
- • Governance challenges

What is a Data Lakehouse?

The Best of Both Worlds



Key Lakehouse Principles

| Principle | Description | Benefit |
|--------------------|--------------------------|--------------------|
| Open Formats | Data in Parquet, Delta | No vendor lock-in |
| ACID Transactions | Reliable updates | Data consistency |
| Schema Enforcement | Data quality at write | Prevent bad data |
| Time Travel | Access previous versions | Auditing, rollback |
| Unified Access | Same data for all | No data silos |
| Direct Access | Query in place | No copying needed |

Delta Lake: The Foundation

What is Delta Lake?

An **open-source storage layer** that brings reliability to Data Lakes.

How it works:

- Stores data as **Parquet files**
- Adds a **transaction log**
- Enables ACID on object storage

Table Structure:

```
delta_table/  
+-- _delta_log/  
|   --- 000...00.json  
|   --- 000...01.json  
|   --- 000...02.json  
+-- part-00000.parquet  
+-- part-00001.parquet  
+-- part-00002.parquet
```

Delta Lake Features

| Feature | Description | Example |
|--------------------|-----------------------------|-----------------------------|
| ACID Transactions | Atomic, reliable operations | Multiple concurrent writers |
| Time Travel | Query historical data | VERSION AS OF 5 |
| Schema Evolution | Add columns safely | ALTER TABLE ADD COL |
| Schema Enforcement | Reject invalid data | Wrong types fail |
| Audit History | Track all changes | DESCRIBE HISTORY |
| MERGE (Upserts) | Update/Insert combined | CDC processing |

Time Travel in Action

Query Any Version of Your Data


Query by Version:

```
-- Current version  
SELECT * FROM sales;  
  
-- Specific version  
SELECT * FROM sales  
VERSION AS OF 5;
```

Query by Timestamp:

```
-- Point in time  
SELECT * FROM sales  
TIMESTAMP AS OF  
'2024-01-01';  
  
-- Restore data  
RESTORE TABLE sales  
TO VERSION AS OF 5;
```

Complete Lakehouse Architecture



Data Ingestion Layer

Auto Loader

- Incremental file ingestion
- Schema inference
- Exactly-once semantics

Streaming

- Kafka integration
- Event Hubs
- Kinesis support

CDC

- Database changes
- Real-time sync
- MERGE operations

Processing & Governance Layers

Processing Layer

- **Apache Spark:** Core engine
- **Photon Engine:** C++ vectorized
(2-8x faster)
- **Delta Live Tables:** Declarative ETL

Governance Layer

- **Unity Catalog:** Centralized governance
- **Data Lineage:** Track data flow
- **Access Control:** Row/column security

Ensuring Data Reliability

| Property | Definition | Example |
|-------------|-------------------------|------------------------------|
| Atomicity | All succeed or all fail | Debit AND credit both happen |
| Consistency | Data always valid | Balance never negative |
| Isolation | No interference | Two concurrent updates |
| Durability | Survives failures | Power outage safe |

How Delta Lake Achieves ACID

1. Optimistic Concurrency


- Writers don't block
- Conflicts at commit
- Automatic retry

2. Write-Ahead Log

- Log before data
- Enables recovery

3. Atomic Commits

- New log entry = complete
- Incomplete writes ignored



Medallion Architecture

Bronze → Silver → Gold



Bronze Layer - Raw Data

Characteristics:

- ▷ Raw data “as-is” from source
- ▷ No transformations
- ▷ Append-only writes
- ▷ Full history preserved
- ▷ Schema-on-read

Example:

- Raw JSON from IoT sensors
- CSV files from vendors
- API response dumps

Quality: Low - may have errors, duplicates

Silver Layer - Cleansed Data

Characteristics:

- ▷ Cleaned and validated
- ▷ Deduplication applied
- ▷ Data types enforced
- ▷ Joined with reference data
- ▷ Schema enforced

Example:

- Parsed sensor readings
- Typed and validated
- Deduplicated records

Quality: Medium - business rules applied

Gold Layer - Business Ready

Characteristics:

- ▷ Business-level aggregations
- ▷ Denormalized for performance
- ▷ Optimized for use cases
- ▷ Multiple tables per domain

Example:

- Hourly averages by location
- Feature tables for ML
- KPI dashboards

Quality: High - analytics ready

Key Takeaways

- **Lakehouse** = Data Lake flexibility + Data Warehouse reliability
- **Delta Lake** provides ACID transactions on object storage
- **Time Travel** enables auditing and rollback capabilities
- **Medallion Architecture:** Bronze → Silver → Gold
- **Unity Catalog** provides centralized governance

One platform for all data workloads!

Thank You!

Questions?

Connect with me:

linkedin.com/in/yashkavaiya

Gen AI Guru