

# Delta Lake Advanced

Time Travel, MERGE, OPTIMIZE & VACUUM

Databricks 14-Days AI Challenge

Day 5 - Advanced Delta Lake Features

January 14, 2026

# Agenda

- **Delta Lake Architecture**
  - ▷ Transaction Log & Parquet Files
  - ▷ ACID Properties
- **Time Travel**
  - ▷ Version History Queries
  - ▷ Restore Operations
- **MERGE Operations**
  - ▷ Upsert Patterns
  - ▷ SCD Type 2
- **OPTIMIZE & ZORDER**
  - ▷ File Compaction
  - ▷ Data Skipping
- **VACUUM for Cleanup**
  - ▷ Stale File Removal
  - ▷ Retention Configuration
- **Best Practices**
  - ▷ Performance Tips
  - ▷ Common Pitfalls


# Delta Lake Architecture Overview

## What is Delta Lake?

An open-source storage layer that brings **ACID transactions** to Apache Spark and big data workloads.

## Core Components:

- ▷ **Transaction Log (\_delta\_log)**
  - ▷ JSON files recording every change
  - ▷ Enables ACID & time travel
- ▷ **Parquet Data Files**
  - ▷ Efficient columnar storage
  - ▷ Optimized compression



# ACID Properties in Delta Lake

## Atomicity

Each transaction completely succeeds or completely fails. Partial changes are never visible.

## Isolation

Concurrent transactions don't interfere. Readers see consistent snapshots while writers make changes.

## Consistency

Table always moves from one valid state to another. Schema enforcement ensures data integrity.

## Durability

Once committed, changes are permanent and survive system failures.



# Time Travel - Query Historical Data

## Three Methods to Access Historical Versions:

### Method 1: Version Number

```
SELECT * FROM my_table  
VERSION AS OF 5;  
  
-- PySpark  
df = spark.read.format("delta")  
    .option("versionAsOf", 5)  
    .load("/path/to/table")
```

### Method 3: Version Shorthand

```
-- Version shorthand syntax  
SELECT * FROM my_table VERSION AS OF  
      5;  
  
-- Timestamp shorthand  
SELECT * FROM my_table  
TIMESTAMP AS OF '20240115103000';
```

### Method 2: Timestamp

```
SELECT * FROM my_table  
TIMESTAMP AS OF '2024-01-15  
    10:30:00';
```

### View Table History:

```
DESCRIBE HISTORY my_table;  
DESCRIBE HISTORY my_table LIMIT 10;
```

# Time Travel - Restore Operations

## Restoring Previous Versions:

```
-- Restore to a specific version  
RESTORE TABLE my_table TO VERSION AS OF 10;  
  
-- Restore to a specific timestamp  
RESTORE TABLE my_table TO TIMESTAMP AS OF '2024-01-15 10:30:00';
```

## Key Use Cases:

- ▷ **Auditing & Compliance**
- ▷ **Data Recovery**
- ▷ **Reproducibility (ML)**
- ▷ **Debugging Issues**
- ▷ **Rollback Pipelines**

## Retention Configuration:

```
-- Log retention (default 30 days)  
ALTER TABLE my_table SET  
TBLPROPERTIES (  
    'delta.logRetentionDuration'  
    = '7 days');  
  
-- File retention (default 7 days)  
ALTER TABLE my_table SET  
TBLPROPERTIES (  
    'delta.  
        deletedFileRetentionDuration',
```

# MERGE Operations (Upserts)

Combine **INSERT**, **UPDATE**, **DELETE** in a single atomic transaction

```
MERGE INTO target_table AS target
USING source_table AS source
ON target.id = source.id
WHEN MATCHED THEN
    UPDATE SET
        target.name = source.name,
        target.value = source.value,
        target.updated_at = current_timestamp()
WHEN NOT MATCHED THEN
    INSERT (id, name, value, created_at)
    VALUES (source.id, source.name, source.value, current_timestamp());
```



# MERGE Clause Types

## WHEN MATCHED:

Update or delete matching rows

```
-- Update matching rows
WHEN MATCHED THEN
    UPDATE SET target.col1 = source.
        col1

-- With condition
WHEN MATCHED AND source.is_deleted =
    true
THEN DELETE
WHEN MATCHED AND source.is_deleted =
    false
THEN UPDATE SET target.col1 = source
    .col1
```

## WHEN NOT MATCHED:

Insert new rows

```
-- Insert new rows
WHEN NOT MATCHED THEN
    INSERT (id, name, value)
    VALUES (source.id, source.name,
            source.value)
```

## WHEN NOT MATCHED BY SOURCE:

Handle orphaned records

```
-- Delete orphaned records
WHEN NOT MATCHED BY SOURCE THEN
    DELETE
```

# MERGE with PySpark

```
from delta.tables import DeltaTable
from pyspark.sql.functions import current_timestamp

# Load the target Delta table
deltaTable = DeltaTable.forPath(spark, "/path/to/customers")

# Prepare source DataFrame
updates_df = spark.read.parquet("/path/to/updates")


# Perform MERGE
deltaTable.alias("target").merge(
    updates_df.alias("source"),
    "target.customer_id = source.customer_id"
).whenMatchedUpdate(
    condition="source.operation = 'UPDATE'",
    set={"name": "source.name", "email": "source.email",
         "modified_date": current_timestamp()}
).whenMatchedDelete(
    condition="source.operation = 'DELETE'"
).whenNotMatchedInsert(
    condition="source.operation = 'INSERT'",
    values={"customer_id": "source.customer_id", "name": "source.name",
            "created_date": current_timestamp()})
```

# The Small File Problem


## Why Small Files are Problematic:

- ▷ **Read Overhead**
  - ▷ Each file requires metadata operations
  - ▷ 1000 Å— 1MB slower than 10 Å— 100MB
- ▷ **Memory Pressure**
  - ▷ Spark tracks all files in memory
- ▷ **Cloud API Costs**
  - ▷ More LIST and GET operations
  - ▷ Object stores charge per API call

Before OPTIMIZE



After OPTIMIZE



# OPTIMIZE Operation

Compacts small files into larger ones (target: 1GB)

## SQL Syntax:

```
-- Basic OPTIMIZE  
OPTIMIZE my_table;  
  
-- Specific partitions  
OPTIMIZE my_table  
WHERE date >= '2024-01-01';
```

## How OPTIMIZE Works:

1. Identifies small files
2. Groups files (respects partitions)
3. Reads and rewrites data
4. Updates transaction log atomically
5. Original files marked for deletion

## PySpark:

```
from delta.tables import DeltaTable  
deltaTable = DeltaTable.forPath(  
    spark, "/path/to/table")  
deltaTable.optimize()  
    .executeCompaction()
```

**Note:** Original files removed by VACUUM

# ZORDER Optimization

Co-locates related data in the same files for better data skipping

```
-- ZORDER on specific columns
OPTIMIZE my_table
ZORDER BY (customer_id);

-- Multiple columns
OPTIMIZE my_table
ZORDER BY (region, customer_id);

-- Combined with partition filter
OPTIMIZE my_table
WHERE date >= '2024-01-01'
ZORDER BY (customer_id, product_id);
```

## Good ZORDER Candidates:

- ✓ High-cardinality filter columns
- ✓ Columns used in JOINs
- ✓ WHERE clause columns

## Poor Candidates:

- ✗ Low-cardinality (use partition)
- ✗ Rarely queried columns
- ✗ Already partitioned columns

# Partitioning vs ZORDER

Characteristic	Use Partitioning	Use ZORDER
Low cardinality ( $< 1000$ )	✓ Yes	✗ No
High cardinality ( $> 1000$ )	✗ No	✓ Yes
Used in every query	✓ Yes	Less important
Used in some queries	✗ No	✓ Yes
Equality predicates only	Yes (if low card)	Yes (if high card)
Range predicates	✗ No	✓ Yes

**Column Order Matters:** First column has strongest locality

# Auto Optimize Features

## Auto Compaction:

Automatically runs OPTIMIZE after writes

```
-- Enable at table level
ALTER TABLE my_table SET
    TBLPROPERTIES (
        'delta.autoOptimize.autoCompact'
        = 'true');

-- Enable for all tables in session
SET spark.databricks.delta
    .autoCompact.enabled = true;
```

## Optimized Writes:

Coalesces small files during writes

```
ALTER TABLE my_table SET
    TBLPROPERTIES (
        'delta.autoOptimize.optimizeWrite'
        = 'true');
```

Feature	When to Use
autoCompact	Streaming
optimizeWrite	Frequent writes
Manual	Batch maintenance

# VACUUM for Cleanup

Removes data files no longer referenced by the transaction log

## Why Files Become Stale:

1. New file written with updates
2. Old file marked as “removed”
3. Old file still exists (for time travel)
4. **VACUUM cleans them up**

## SQL Syntax:

```
-- Default 7-day retention
VACUUM my_table;

-- Custom retention
VACUUM my_table RETAIN 168 HOURS;

-- Dry run first!
VACUUM my_table DRY RUN;
```

## PySpark:

```
from delta.tables import DeltaTable

deltaTable = DeltaTable.forPath(
    spark, "/path/to/table")

# Default retention
deltaTable.vacuum()

# Custom retention (hours)
deltaTable.vacuum(168)
```

**Warning:** VACUUM affects time travel!

# VACUUM Retention & Safety

## Recommended Retention Periods:

Scenario	Retention	Reasoning
Production tables	7-30 days	Balance cost & recovery
Development tables	1-7 days	Lower cost, less recovery
Audit-required tables	30-365 days	Compliance requirements
High-frequency updates	7 days min	Protect long queries

## Align VACUUM with Time Travel:

```
-- If you need 30 days of time travel
ALTER TABLE my_table SET TBLPROPERTIES (
    'delta.logRetentionDuration' = '30 days',
    'delta.deletedFileRetentionDuration' = '30 days');

-- Then VACUUM with matching retention
VACUUM my_table RETAIN 720 HOURS; -- 30 days
```

# OPTIMIZE vs VACUUM Comparison

Aspect	OPTIMIZE	VACUUM
Purpose	Improve query performance	Reduce storage usage
Creates new files	✓ Yes	✗ No
Deletes files	✗ No (marks stale)	✓ Yes
Affects time travel	✗ No	✓ Yes
Storage impact	Temporary increase	Decrease
When to run	After many writes	After OPTIMIZE
Frequency	Daily or weekly	Weekly or monthly

Typical Workflow: Write Data → OPTIMIZE → VACUUM

# Best Practices & Performance Tips

## Table Design:

- ▶ Partition by low-cardinality columns
- ▶ Aim for 1GB+ partitions
- ▶ Avoid  $\geq 10,000$  partitions
- ▶ ZORDER high-cardinality columns

## MERGE Optimization:

- ▶ Pre-filter source data
- ▶ Include partition columns in join
- ▶ Batch small merges

## OPTIMIZE Schedule:

- ▶ Streaming: Every 1-4 hours
- ▶ Micro-batch: Daily
- ▶ Batch: After each load

## VACUUM Guidelines:

- ▶ Never  $\geq 7$  days without understanding
- ▶ Always DRY RUN first
- ▶ Run during low-traffic periods
- ▶ Schedule after OPTIMIZE

# Common Pitfalls to Avoid

Pitfall	Problem	Solution
Over-partitioning	Too many small files	Fewer partitions, use ZORDER
ZORDER on partition cols	Redundant, no benefit	ZORDER non-partition columns
Aggressive VACUUM	Lose time travel	Match retention to needs
No OPTIMIZE schedule	Small file problem	Automate with Auto Optimize
MERGE without filters	Full table rewrite	Pre-filter, partition pruning
Ignoring file stats	Poor data skipping	ZORDER filtered columns

# Quick Reference Card

## Time Travel Commands:

```
SELECT * FROM t VERSION AS OF 5;
SELECT * FROM t TIMESTAMP AS OF
  '2024-01-15';
DESCRIBE HISTORY t;
RESTORE TABLE t TO VERSION AS OF 5;
```

## MERGE Pattern:

```
MERGE INTO target USING source
ON condition
WHEN MATCHED THEN UPDATE SET ...
WHEN NOT MATCHED THEN INSERT ...
```

## OPTIMIZE & VACUUM:

```
OPTIMIZE t;
OPTIMIZE t ZORDER BY (col1, col2);
OPTIMIZE t WHERE partition_col = 'v
  ';
VACUUM t;
VACUUM t RETAIN 168 HOURS;
VACUUM t DRY RUN;
```

## Key Properties:

```
'delta.logRetentionDuration',
'delta.deletedFileRetentionDuration',
'delta.autoOptimize.optimizeWrite',
'delta.autoOptimize.autoCompact'
```

# Thank You!

Questions?

 [linkedin.com/in/yashkavaiya](https://www.linkedin.com/in/yashkavaiya)

 [easy-ai-labs.lovable.app](https://easy-ai-labs.lovable.app)

 Gen AI Guru